

Linear Regression

- ◆ Linear regression predicts a real value output based on the input value.
 - Simple linear regression only has one input/variable and we try to fit a line to best describe the output and input relationship
 - Multiple linear regression has multiple variables as input and we try to fit a hyperplane instead.
- ◆ An example of simple linear regression is salary prediction, with the salary as output, and years of working experience as input variable.
- ◆ The algorithm can be used for simple linear regression or multiple linear regression, and in this notebook, it will be implemented for simple linear regression using [sklearn linear regression algorithm](#)

Import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.linear_model import LinearRegression
```

Read data

```
In [2]: df = pd.read_csv('Salary_Data.csv')
# shuffle data
df = df.sample(frac=1, random_state=42).reset_index(drop=True)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	YearsExperience	Salary
0	9.6	112635.0
1	4.9	67938.0
2	8.2	113812.0
3	5.3	83088.0
4	3.2	64445.0

```
In [4]: # visualize data
fig = plt.figure(figsize=(8,6))
plt.scatter(df.iloc[:,0], df.iloc[:,1], c='b')
plt.title('Salary vs. Years of Experience', fontsize=14)
plt.xlabel('Years of Experience', fontsize=12)
plt.ylabel('Salary', fontsize=12)
plt.show()
```



Split data to train and test dataset

```
In [5]: x = np.array(df.iloc[:,0:1])  
y = np.array(df.iloc[:,1])
```

```
In [6]: train_frac = 0.7  
train_size = int(df.shape[0] * train_frac)  
x_train = x[0:train_size,:]  
y_train = y[0:train_size]  
x_test = x[train_size:,:]  
y_test = y[train_size:]
```

```
In [7]: print ('training size is {}\ntest size is {}'.format(x_train.shape[0], x_test.shape[0]))
```

```
training size is 21  
test size is 9
```

Initialize the model

```
In [8]: lr = LinearRegression()  
lr.fit(x_train, y_train)  
train_accuracy = lr.score(x_train, y_train)*100
```

Predict the test data output

```
In [9]: y_pred_train = lr.predict(x_train)  
y_pred_test = lr.predict(x_test)
```

```
coef = lr.coef_
intercept = lr.intercept_
test_accuracy = lr.score(x_test, y_test)*100
```

In [10]:

```
# print the test data, predicted test data and the error
np.set_printoptions(formatter={'float': lambda x: "{0:0.01f}".format(x)})
print('y_test\n {} \n\n y_pred_test\n {} \n\n test accuracy (%) {} \n\n coef {} intercept {}'
      .format(y_test, y_pred_test, test_accuracy, coef, intercept))
```

```
y_test
[116969.0  81363.0 121872.0  91738.0  54445.0  63218.0  61111.0  93940.0  60150.0]
```

```
y_pred_test
[115224.5  81135.1 124693.7  89657.4  55568.1  62196.6  67878.1  82082.0  53674.2]
```

```
test accuracy (%) 95.1483199421986
```

```
coef [9469.3] intercept 25266.40502593999
```

Visualize fitted results

In [11]:

```
fig = plt.figure(figsize=(8,6))
plt.scatter(df.iloc[:,0], df.iloc[:,1], c='b')
plt.plot(x_train, y_pred_train, c='r', lw=2)
plt.title('', fontsize=14)
plt.xlabel('Years of Experience', fontsize=12)
plt.ylabel('Salary', fontsize=12)
plt.legend(['fitted line using sklearn LinearRegression'], fontsize=12)
plt.show()
```

