

Logistic regression

- ◆ Logistic regression is an algorithm that can help us answering yes or no questions by predicting the probability something happening or not happening.
- ◆ Logistic regression is a classification algorithm despite its name has regression in it. It predicts two binary dependent output either as 0 or 1 based on the input variables.
- ◆ The regression in its name means we are using the same algorithm as used in linear regression. The difference is that the output is mapped using a logistic/sigmoid function so that it will be in a range between 0 and 1. Any output value < 0.5 will be classified as 0 and any value ≥ 0.5 will be classified as 1.
- ◆ In this notebook, the logistic regression algorithm is coded step by step following:

Logistic regression

Hypothesis: $h_{\theta}(x) = \sigma(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$

parameters/weights: $\theta = [\theta_0, \theta_1, \theta_2 \dots \theta_n]$

Cost function: $J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^i \log(h_{\theta}(x^i)) + (1 - y^i) \log(1 - h_{\theta}(x^i))]$

Find the parameters that'll minimize $J(\theta)$: Gradient Descent

△ Start with Random θ

△ update the weights until reach minimum of $J(\theta)$

$$\theta_j = \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^i) - y^i) x_j^i$$

Note: All the weights should be updated simultaneously

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Import libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

Coding the algorithm

```
In [2]: class LogisticRegression():
def __init__(self, x, y):
```

```

x0 = np.ones((x.shape[0],1)) # add the bias term
x = np.append(x0, x, axis=1) # append the bias to the input variable
self.x = x
self.y = y
self.m = x.shape[0]
self.n = x.shape[1]
self.theta = np.random.randn(self.n) # initialize the weights with random num

def sigmoid(self, z):
    return 1/(1+np.exp(-z))

def CostFunction(self):
    self.h = self.sigmoid(np.matmul(self.x, self.theta)) # calculate the hypotheses
    self.J = (-1/self.m) * np.sum( self.y*np.log(self.h) + (1-self.y)*np.log(1 -
#         print (self.J)
    return self.h, self.J

def GradientDescent(self, epoch=10, alpha=0.01):
    self.cost_history = []
    self.theta_history = []
    for i in range(epoch):
        h, J = self.CostFunction()
        self.cost_history.append(J)
        self.theta_history.append(self.theta)
        self.theta = self.theta - alpha / self.m * np.dot(self.x.T, h-self.y)
    pass

def Theta(self):
    return self.theta

def CostHistory(self):
    return self.cost_history

def ThetaHistory(self):
    return self.theta_history

def predict(self, x_test, y_test):
    x0 = np.ones((x_test.shape[0],1))
    x_test = np.append(x0, x_test, axis=1)
    self.y_pred = self.sigmoid(np.matmul(x_test, self.theta))
    self.y_pred[self.y_pred>=0.5] = 1
    self.y_pred[self.y_pred<0.5] = 0
    self.test_accuracy = (self.y_pred == y_test).sum()/ len(y_test)*100
    return self.y_pred.astype(int), self.test_accuracy

```

```

In [3]: # feature scaling is not used in this notebook, coded for fun.
class FeatureScaler():
    def __init__(self, x):
        self.x = x.copy().astype(float)
    def FitTransform(self):
        n = self.x.shape[1]
        for i in range(n):
            xi = self.x[:,i]
            xi_scaled = (xi - xi.mean()) / (xi.max() - xi.min())

```

```
self.x[:,i] = xi_scaled
return self.x
```

Simulating data

The data used in this notebook are simulated using `np.random.multivariate_normal` to create two seperable features. Credit to [git repo](#)


In [4]:

```
np.random.seed(42)
num_observations = 1000

x1 = np.random.multivariate_normal([0, 0], [[1, .75],[.75, 1]], num_observations)
x2 = np.random.multivariate_normal([1, 4], [[1, .75],[.75, 1]], num_observations)
features = np.vstack((x1, x2)).astype(np.float32)
labels = np.hstack((np.zeros(num_observations),
                    np.ones(num_observations)))

#Displaying the output

plt.figure(figsize=(12,8))
plt.scatter(features[:, 0], features[:, 1], c = labels, alpha = .4)
plt.show()
```



Split data to train and test dataset

In [5]:

```
df = pd.DataFrame()
```

```
In [6]: df['feature_1'] = features[:,0]
df['feature_2'] = features[:,1]
df['label'] = labels.astype(int)
```

```
In [7]: df.head()
```

```
Out[7]:
```

	feature_1	feature_2	label
0	-0.415750	-0.513517	0
1	-1.144330	-0.067385	0
2	0.301810	0.136251	0
3	-1.748547	-1.205889	0
4	0.247329	0.630977	0

```
In [8]: df = df.sample(frac=1, random_state=42).reset_index(drop=True)
```

```
In [9]: train_frac = 0.7
train_size = int(df.shape[0] * train_frac)
x = np.array(df.iloc[:,0:2])
y = np.array(df.iloc[:,2])
x_train = x[0:train_size,:]
y_train = y[0:train_size]
x_test = x[train_size:,:]
y_test = y[train_size:]
```

```
In [10]: print ('training size is {}\ntest size is {}'.format(x_train.shape[0], x_test.shape[0])

training size is 1400
test size is 600
```

Initialize the model

```
In [11]: clf = LogisticRegression(x_train, y_train)
```

Predict using gradient descent method

```
In [12]: epoch = 1000
alpha = 0.5
clf.GradientDescent(epoch, alpha)
y_pred_train, train_error = clf.predict(x_train,y_train)
y_pred_test, accuracy = clf.predict(x_test,y_test)
theta = clf.Theta()
```

```
In [13]: # print the test data, predicted test data and the error
print('y_test\n {}\n\ny_pred_test\n {}\n\ntest accuracy (%) {}\n\ntheta {}')
```

```
.format(y_test[0:20], y_pred_test[0:20], accuracy, theta))
```

```
y_test  
[0 0 0 1 1 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1]
```

```
y_pred_test  
[0 0 0 1 1 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1]
```


```
test accuracy (%) 99.5
```

```
theta [-6.51769471 -1.92665152  3.77311985]
```

Visulize cost vs epochs

In [14]:


```
cost = clf.CostHistory()  
fig = plt.figure(figsize=(8,6))  
plt.plot(range(epoch), cost, c='b')  
plt.title('Cost vs. No. of Epochs', fontsize=14)  
plt.xlabel('No. of Epochs', fontsize=12)  
plt.ylabel('Cost', fontsize=12)  
plt.show()
```



Visulize fitted results using both method


In [15]:

```
fig = plt.figure(figsize=(8,6))  
plt.scatter(x_test[:,0], x_test[:,1], c=y_pred_test, alpha = .5)  
plt.title('', fontsize=14)  
plt.xlabel('Years of Experience', fontsize=12)  
plt.ylabel('Salary', fontsize=12)  
plt.show()
```



In [16]:

```
fig = plt.figure(figsize=(8,6))
plt.scatter(x_test[:,0], x_test[:,1], c=y_test, alpha = .5)
plt.title('', fontsize=14)
plt.xlabel('Years of Experience', fontsize=12)
plt.ylabel('Salary', fontsize=12)
plt.show()
```



In []: