

# Logistic regression

- ◆ Logistic regression is an algorithm that can help us answering yes or no questions by predicting the probability something happening or not happening.
- ◆ Logistic regression is a classification algorithm despite its name has regression in it. It predicts two binary dependent output either as 0 or 1 based on the input variables.
- ◆ The regression in its name means we are using the same algorithm as used in linear regression. The difference is that the output is mapped using a logistic/sigmoid function so that it will be in a range between 0 and 1. Any output value  $< 0.5$  will be classified as 0 and any value  $\geq 0.5$  will be classified as 1.
- ◆ In this notebook, the [sklearn logistic regression algorithm](#) is implemented on a simulated seperable dataset with two features

## Import libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.linear_model import LogisticRegression
```

## Simulating data

The data used in this notebook are simulated using `np.random.multivariate_normal` to create two seperable features. Credit to [git repo](#)


In [2]:

```
np.random.seed(42)
num_observations = 1000

x1 = np.random.multivariate_normal([0, 0], [[1, .75],[.75, 1]], num_observations)
x2 = np.random.multivariate_normal([1, 4], [[1, .75],[.75, 1]], num_observations)
features = np.vstack((x1, x2)).astype(np.float32)
labels = np.hstack((np.zeros(num_observations),
                    np.ones(num_observations)))

#Displaying the output

plt.figure(figsize=(12,8))
plt.scatter(features[:, 0], features[:, 1], c = labels, alpha = .4)
plt.show()
```



## Split data to train and test dataset

```
In [3]: df = pd.DataFrame()
```

```
In [4]: df['feature_1'] = features[:,0]
df['feature_2'] = features[:,1]
df['label'] = labels.astype(int)
```

```
In [5]: df.head()
```

```
Out[5]:
```

	feature_1	feature_2	label
0	-0.415750	-0.513517	0
1	-1.144330	-0.067385	0
2	0.301810	0.136251	0
3	-1.748547	-1.205889	0
4	0.247329	0.630977	0

```
In [6]: df = df.sample(frac=1, random_state=42).reset_index(drop=True)
```

```
In [7]: train_frac = 0.7
```

```

train_size = int(df.shape[0] * train_frac)
x = np.array(df.iloc[:,0:2])
y = np.array(df.iloc[:,2])
x_train = x[0:train_size,:]
y_train = y[0:train_size]
x_test = x[train_size:,:]
y_test = y[train_size:]

```

In [8]: `print ('training size is {}\ntest size is {}'.format(x_train.shape[0], x_test.shape[0])`

```

training size is 1400
test size is 600

```

## Initialize the model

In [9]:

```

clf = LogisticRegression(random_state=42)
clf.fit(x_train, y_train)
coef = clf.coef_
intercept = clf.intercept_
train_accuracy = clf.score(x_train, y_train)
print ('training accuracy {}'.format(train_accuracy))

```

```

training accuracy 0.9935714285714285

```

## Predict using gradient descent method

In [10]:

```

y_pred_train = clf.predict(x_train)
y_pred_test = clf.predict(x_test)
test_accuracy = clf.score(x_test, y_test)

```

In [11]:

```

# print the test data, predicted test data and the error
print('y_test\n {}\n\ny_pred_test\n {}\n\ntest accuracy {}\n\ncoef {} intercept {}'.format(y_test[0:20], y_pred_test[0:20], test_accuracy, coef, intercept))

```

```

y_test
[0 0 0 1 1 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1]

```

```

y_pred_test
[0 0 0 1 1 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1]

```

```

test accuracy 0.995

```

```

coef [[-2.22854218  4.55550901]] intercept [-8.20091535]

```


## Visualize fitted results using both method

In [12]:


```

fig = plt.figure(figsize=(8,6))
plt.scatter(x_test[:,0], x_test[:,1], c=y_pred_test)
plt.title('', fontsize=14)
plt.xlabel('Years of Experience', fontsize=12)
plt.ylabel('Salary', fontsize=12)
plt.show()

```



```
In [13]: fig = plt.figure(figsize=(8,6))
plt.scatter(x_test[:,0], x_test[:,1], c=y_test)
plt.title('', fontsize=14)
plt.xlabel('Years of Experience', fontsize=12)
plt.ylabel('Salary', fontsize=12)
plt.show()
```



```
In [ ]:
```